

Multi-layer Perceptron

Chang Lu

1 Model

Let $\mathbf{x} \in \mathbb{R}^{s \times n}$ be the inputs, where n is the number of samples; and s is the dimension of each feature vector. We first define a hidden layer t in multi-layer perceptron (MLP):

$$\mathbf{z}^{(t)} = \mathbf{w}^{(t)} \mathbf{h}^{(t-1)} + \mathbf{b}^{(t)}, \quad (1)$$

$$\mathbf{h}^{(t)} = g^{(t)}(\mathbf{z}^{(t)}). \quad (2)$$

Here, $\mathbf{h}^{(t)} \in \mathbb{R}^{d^{(t)} \times n}$ is the output of the t -th layer. $\mathbf{h}^{(0)} = \mathbf{x}$. $\mathbf{w}^{(t)} \in \mathbb{R}^{d^{(t)} \times d^{(t-1)}}$ is the weight to map the output of the $(t-1)$ -th layer into the intermediate output \mathbf{z} . $\mathbf{b}^{(t)} \in \mathbb{R}^{d^{(t)} \times 1}$. $g^{(t)}$ is an activation function such as ReLU, sigmoid, or tanh. Typically, $g^{(t)}$ is the same among hidden layers in MLP, except the output layer. We denote the activation function in hidden layers as g .

In the output layer, the activation function and dimension of weights depend on specific tasks and labels. Let C be the output dimension and f be the activation function, the output layer is defined as follows:

$$\mathbf{z}^{(T)} = \mathbf{w}^{(T)} \mathbf{h}^{(T-1)} + \mathbf{b}^{(T)}, \quad (3)$$

$$\hat{\mathbf{y}} = f(\mathbf{z}^{(T)}). \quad (4)$$

Here, $\mathbf{w}^{(T)} \in \mathbb{R}^{C \times d^{(T-1)}}$. Traditionally, there is an argument about counting the layers in MLP. In our project, we count one layer by the weight w . Figure 1 shows an example of a 2-layer MLP. It contains two hidden weight variables. We call it 2-layer MLP (input neurons, hidden neurons, and output neurons).




Figure 1: A 2-layer MLP

2 Objective Function

The objective function depends on various tasks. We use a multi-class classification as an example. The objective function is selected as multi-class cross-entropy loss. Let C be the output dimension, i.e., number of categories and $\mathbf{y} \in \{0, 1\}^{C \times n}$ be a one-hot vector for the ground-truth, the prediction of MLP is $\hat{\mathbf{y}} \in \mathbb{R}^{C \times n}$ is a probability distribution of each entry from $\mathbf{z}^{(T)}$, calculated by a softmax activation function:

$$\hat{y}_c = \frac{e^{\mathbf{z}_c^{(T)}}}{\sum_{k=1}^C e^{\mathbf{z}_k^{(T)}}}. \quad (5)$$

For an input sample \mathbf{x}_i , the multi-class cross-entropy loss $L(\mathbf{x}_i, \mathbf{y}_i \mid \mathbf{w}^{(1)}, \dots, \mathbf{w}^{(T)})$ is:

$$L(\mathbf{x}_i, \mathbf{y}_i \mid \mathbf{w}^{(1)}, \dots, \mathbf{w}^{(T)}) = - \sum_{c=1}^C \mathbf{y}_{c,i} \log \hat{\mathbf{y}}_{c,i} \quad (6)$$

$$\begin{aligned} &= - \sum_{c=1}^C \mathbf{y}_{c,i} \left(\mathbf{z}_{c,i}^{(T)} - \log \sum_{k=1}^C e^{\mathbf{z}_{i,k}^{(T)}} \right) \\ &= \log \sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}} - \sum_{c=1}^C \mathbf{y}_{c,i} \mathbf{z}_{c,i}^{(T)}. \end{aligned} \quad (7)$$

For all input samples \mathbf{x} , the total loss is an average of the losses for all samples:

$$\begin{aligned} L(\mathbf{x}, \mathbf{y} \mid \mathbf{w}^{(1)}, \dots, \mathbf{w}^{(T)}) &= -\frac{1}{n} \sum_{i=1}^n \sum_{c=1}^C \mathbf{y}_{c,i} \log \hat{\mathbf{y}}_{c,i} \\ &= \frac{1}{n} \sum_{i=1}^n \left(\log \sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}} - \sum_{c=1}^C \mathbf{y}_{c,i} \mathbf{z}_{c,i}^{(T)} \right). \end{aligned} \quad (8)$$

2.1 A Trick for Stable Softmax and Log-Sum-Exp

When calculating e^x for softmax function in Equation (5) and Equation (7), x can be large so that there may be a numerical problem. To alleviate this problem, we can take advantage a property of softmax:

$$\begin{aligned} \text{softmax}(x_c \mid x_1, \dots, x_C) &= \frac{e^{x_c}}{\sum_{k=1}^C e^{x_k}} = \frac{e^{-\lambda} \cdot e^{x_c}}{e^{-\lambda} \cdot \sum_{k=1}^C e^{x_k}} = \frac{e^{x_c - \lambda}}{\sum_{k=1}^C e^{x_k - \lambda}} \\ &= \text{softmax}(x_c - \lambda \mid x_1 - \lambda, \dots, x_C - \lambda). \end{aligned} \quad (9)$$

Similarly, we can get

$$\log \sum_{c=1}^C e^{x_c} = \log \left(e^\lambda \cdot e^{-\lambda} \cdot \sum_{c=1}^C e^{x_c} \right) = \lambda + \log \sum_{c=1}^C e^{x_c - \lambda}. \quad (10)$$

Let $\lambda = \max \{x_1, x_2, \dots, x_C\}$, for $\forall c \geq 1$ and $c \leq C$, we have $x_c - \lambda \leq 0 \Rightarrow e^{x_c - \lambda} \leq 1$. In this way, the softmax and log-sum-exp operations can be numerically stable.

3 Back-propagation

We use the multi-class classification for predictions and ReLU activation for hidden layers. The trainable variables are $\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(T)}$ and $\mathbf{b}^{(1)}, \dots, \mathbf{b}^{(T)}$. We need to calculate gradients for all of these variables. For $\mathbf{w}^{(t)}$, we need to apply the chain rule:

$$\frac{\partial L}{\partial \mathbf{w}^{(t)}} = \frac{\partial L}{\partial \mathbf{z}^{(t)}} \frac{\partial \mathbf{z}^{(t)}}{\partial \mathbf{w}^{(t)}} \in \mathbb{R}^{d^{(t-1)} \times d^{(t)}}.$$

For the first term, it is not easy to directly get the gradient of $\mathbf{z}^{(t)}$. For the second term, it is a matrix-by-matrix derivative, we cannot directly calculate either. Therefore, we first consider the loss of one sample: $L_i = L(\mathbf{x}_i, \mathbf{y}_i)$. Then we seek to pass gradients from the $(t+1)$ -th layer w.r.t. the k -th row of \mathbf{w}^t . Let

$$\begin{aligned} \phi(t, i) &= \frac{\partial L_i}{\partial \mathbf{z}_{:,i}^{(t)}} \in \mathbb{R}^{1 \times d^{(t)}}, \\ \psi(t, i, k) &= \frac{\partial L_i}{\partial \mathbf{w}_{k,:}^{(t)}} \in \mathbb{R}^{1 \times d^{(t-1)}}, \end{aligned}$$

we have

$$\begin{aligned}\phi(t, i) &= \underbrace{\phi(t+1, i)}_{1 \times d^{(t+1)}} \underbrace{\frac{\partial \mathbf{z}_{:,i}^{(t+1)}}{\partial \mathbf{z}_{:,i}^{(t)}}}_{d^{(t+1)} \times d^{(t)}} , \\ \psi(t, i, k) &= \underbrace{\phi(t, i)}_{1 \times d^{(t)}} \underbrace{\frac{\partial \mathbf{z}_{:,i}^{(t)}}{\partial \mathbf{w}_{k,:}^{(t)}}}_{d^{(t)} \times d^{(t-1)}} .\end{aligned}$$

In this way, we only need to calculate $\phi(T, i)$, $\frac{\partial \mathbf{z}_{:,i}^{(t+1)}}{\partial \mathbf{z}_{:,i}^{(t)}}$, and $\frac{\partial \mathbf{z}_{:,i}^{(t)}}{\partial \mathbf{w}_{k,:}^{(t)}}$, which is easier and intuitive. Then a recursive way can be applied to calculate each $\mathbf{w}^{(t)}$.

$$\begin{aligned}\phi(T, i) &= \frac{\partial}{\partial \mathbf{z}_{:,i}^{(T)}} \left(\log \sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}} - \sum_{c=1}^C \mathbf{y}_{c,i} \mathbf{z}_{c,i}^{(T)} \right) \\ &= \frac{\partial}{\partial \mathbf{z}_{:,i}^{(T)}} \log \sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}} - \frac{\partial}{\partial \mathbf{z}_{:,i}^{(T)}} \sum_{c=1}^C \mathbf{y}_{c,i} \mathbf{z}_{c,i}^{(T)} \\ &= \frac{1}{\sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}}} \cdot \frac{\partial}{\partial \mathbf{z}_{:,i}^{(T)}} \left(\sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}} \right) - \frac{\partial}{\partial \mathbf{z}_{:,i}^{(T)}} \left(\sum_{c=1}^C \mathbf{y}_{c,i} \mathbf{z}_{c,i}^{(T)} \right) \\ &= \frac{1}{\sum_{c=1}^C e^{\mathbf{z}_{c,i}^{(T)}}} \cdot e^{\mathbf{z}_{:,i}^{(T) \top}} - \mathbf{y}_{:,i}^\top \\ &= (\hat{\mathbf{y}}_{:,i} - \mathbf{y}_{:,i})^\top . \\ \frac{\partial \mathbf{z}_{:,i}^{(t+1)}}{\partial \mathbf{z}_{:,i}^{(t)}} &= \frac{\partial \mathbf{z}_{:,i}^{(t+1)}}{\partial \mathbf{h}_{:,i}^{(t)}} \frac{\partial \mathbf{h}_{:,i}^{(t)}}{\partial \mathbf{z}_{:,i}^{(t)}} = \mathbf{w}^{(t+1)} \text{Diag} \left(\text{ReLU}' \left(\mathbf{z}_{:,i}^{(t)} \right) \right) .\end{aligned}$$

Here, $\text{ReLU}'(\cdot)$ is the derivative of ReLU function. We will give its formal definition later. $\text{Diag}(\cdot)$ is a diagonal matrix of which the main diagonal is the given input vector. Therefore, we can derive

$$\begin{aligned}\phi(t, i) &= \phi(t+1, i) \frac{\partial \mathbf{z}_{:,i}^{(t+1)}}{\partial \mathbf{z}_{:,i}^{(t)}} \\ &= \phi(t+1, i) \mathbf{w}^{(t+1)} \text{Diag} \left(\text{ReLU}' \left(\mathbf{z}_{:,i}^{(t)} \right) \right) \\ &= \left(\phi(t+1, i) \mathbf{w}^{(t+1)} \right) \odot \text{ReLU}' \left(\mathbf{z}_{:,i}^{(t)} \right)^\top\end{aligned}$$

Here, \odot denotes element-wise product for two column/row vectors. For $\frac{\partial \mathbf{z}_{:,i}^{(t)}}{\partial \mathbf{w}_{k,:}^{(t)}}$, it is a $d^{(t)} \times d^{(t)}$ matrix. We denote it as \mathbf{X} .

$$\frac{\partial \mathbf{z}_{:,i}^{(t)}}{\partial \mathbf{w}_{k,:}^{(t) \top}} = \mathbf{X},$$

$$\text{where } \mathbf{X}_{j,:} = \begin{cases} \mathbf{h}_{:,i}^{(t-1) \top} & \text{if } j = k, \\ \mathbf{0} & \text{otherwise.} \end{cases}$$

Therefore, we can derive

$$\psi(t, i, k) = \phi(t, i) \mathbf{X} = \phi(t, i)_m \cdot \mathbf{X}_{k,:} = \phi(t, i)_m \cdot \mathbf{h}_{:,i}^{(t-1) \top} .$$

Since $\psi(t, i, k)$ is for the k -th row of $\mathbf{w}^{(t)}$, for the complete $\mathbf{w}^{(t)}$, we have

$$\frac{\partial L_i}{\partial \mathbf{w}^{(t)}} = \psi(t, i) = \left[\psi(t, i, 1)^\top, \psi(t, i, 2)^\top, \dots, \psi(t, i, d^{(t)})^\top \right] = \mathbf{h}_{:,i}^{(t-1)} \phi(t, i).$$

Here, we have already got the gradient for one sample. For all samples, we need to rewrite the $\phi(t, i)$ and $\psi(t, i)$ to a matrix form $\phi(t)$ and $\psi(t)$:

$$\frac{\partial L}{\partial \mathbf{z}^{(T)}} = \phi(T) = (\hat{\mathbf{y}} - \mathbf{y})^\top \in \mathbb{R}^{n \times C}, \quad (11)$$

$$\frac{\partial L}{\partial \mathbf{z}^{(t)}} = \phi(t) = \left(\phi(t+1) \mathbf{w}^{(t+1)} \right) \odot \text{ReLU}' \left(\mathbf{z}^{(t)} \right)^\top \in \mathbb{R}^{n \times d^{(t)}}, \quad (12)$$

$$\frac{\partial L}{\partial \mathbf{w}^{(t)}} = \psi(t) = \frac{1}{n} \sum_{i=1}^n \psi(t, i) = \frac{1}{n} \mathbf{h}^{(t-1)} \phi(t) \in \mathbb{R}^{d^{(t-1)} \times d^{(t)}}. \quad (13)$$

Similarly, the gradient of $\mathbf{b}^{(t)}$ can be calculated as:

$$\frac{\partial L}{\partial \mathbf{b}^{(t)}} = \frac{1}{n} \sum_{i=1}^n \phi(t, i) \in \mathbb{R}^{1 \times d^{(t)}}. \quad (14)$$

Equations (11)-(14) are the final gradients for MLP. For the dimension of matrix derivative, please refer to [Matrix Calculus at Wikipedia](#).

Recap. We may notice that the gradients of the output layer ($t = T$) of MLP is similar to the logistic regression. In fact, the binary classification is a special form of multi-class classification, and logistic regression can be regraded as a single-layer MLP. In practice, the label of binary classification is usually 0 or 1, while the label of multi-class classification is a one-hot vector. This is the reason why we distinguish binary classification from multi-class classification.

After getting the gradients of weights and bias in MLP, there is still a derivative of ReLU activation function to be solved. The ReLU function is defined as follows:

$$\text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0, \\ x & \text{otherwise,} \end{cases} = \max\{0, x\}. \quad (15)$$

It is a convex function that is sub-differential when $x \neq 0$. However, it is undifferentiable at $x = 0$. Therefore, we use subgradient of ReLU and set the gradient as 0 at $x = 0$:

$$\text{ReLU}'(x) = \begin{cases} 0 & \text{if } x \leq 0, \\ 1 & \text{otherwise.} \end{cases} \quad (16)$$